

VE281

Data Structures and Algorithms

Hash Table Size, Rehashing, and Applications of Hashing

Learning Objectives:

- Know how to determine hash table size
- Know why rehashing is needed and how to rehash
- Know amortized analysis
- Know a few typical applications of hashing

Outline

- Hash Table Size and Rehashing
- Applications of Hashing

Determine Hash Table Size

- First, given **performance** requirements, determine the maximum permissible **load factor**.
- Example: we want to design a hash table based on **linear probing** so that **on average**
 - An **unsuccessful** search requires no more than 13 compares.
 - A **successful** search requires no more than 10 compares.

$$U(L) = \frac{1}{2} \left[1 + \left(\frac{1}{1-L} \right)^2 \right] \leq 13 \Rightarrow L \leq \frac{4}{5}$$

$$S(L) = \frac{1}{2} \left[1 + \frac{1}{1-L} \right] \leq 10 \Rightarrow L \leq \frac{18}{19}$$

$$L \leq \frac{4}{5}$$

Determine Hash Table Size

- For a fixed table size, estimate maximum number of items that will be inserted.

- Example: no more than 1000 items.

- For load factor $L = \frac{|S|}{n} \leq \frac{4}{5}$, table size

$$n \geq \frac{5}{4} \cdot 1000 = 1250$$

- Pick n as a **prime** number. For example, $n = 1259$.

However, sometimes there is no limit on the number of items to be inserted.

Rehashing

Motivation

- With more items inserted, the load factor increases. At some point, it will exceed the threshold ($4/5$ in the previous example) determined by the performance requirement.
- For the separate chaining scheme, the hash table becomes inefficient when load factor L is too high.
 - If the size of the hash table is fixed, search performance deteriorates with more items inserted.
- Even worse, for the open addressing scheme, when the hash table becomes full, we **cannot** insert a new item.

Rehashing

- To solve these problems, we need to **rehash**:
 - Create a **larger** table, scan the current table, and then insert items into new table using the new hash function.
 - **Note**: The order is from the beginning to the end of the current table. Not original insertion order.
- We can approximately double the size of the current table.
- **Observation**: The single operation of rehashing is time-consuming. However, it does not occur frequently.
 - How should we justify the time complexity of rehashing?

Amortized Analysis

- **Amortized analysis**: A method of analyzing algorithms that considers the entire sequence of operations of the program.
 - The idea is that while certain operations may be costly, they don't occur frequently; the less costly operations are much more than the costly ones in the long run.
 - Therefore, the cost of those expensive operations is **averaged** over a sequence of operations.
 - In contrast, our previous complexity analysis only considers a single operation, e.g., insert, find, etc.

Amortized Analysis of Rehashing

- Suppose the threshold of the load factor is 0.5. We will double the table size after reaching the threshold.
- Suppose we start from an empty hash table of size $2M$.
- Assume $O(1)$ operation to insert up to M items.
 - Total cost of inserting the first M items: $O(M)$
- For the $(M + 1)$ -th item, create a new hash table of size $4M$.
 - Cost: $O(1)$
- Rehash all M items into the new table. Cost: $O(M)$
- Insert new item. Cost: $O(1)$

Total cost for inserting $M + 1$ items is $2O(M) + 2O(1) = O(M)$.

Amortized Analysis of Rehashing

Total cost for inserting $M + 1$ items is $O(M)$.

- The average cost to insert $M + 1$ items is $O(1)$.
 - Rehashing cost is **amortized** over individual inserts.

Outline

- Hash Table Size and Rehashing
- Applications of Hashing

Application: De-Duplication

- Given: a stream of objects
 - Linear scan through a huge file
 - Or, objects arriving in real time
- Goal: remove duplicates (i.e., keep track of unique objects)
 - E.g., report unique visitors to website
 - Or, avoid duplicates in search result
- Solution: when new object x arrives,
 - Look x in hash table H
 - If not found, insert x into H

Application: 2-SUM Problem

- Given: an unsorted array A of n **distinct** integers. Target sum t .
- Goal: determine whether or not there are two numbers x and y in A with

$$x + y = t$$

1. Naïve solution: exhaustive search of pairs of number
 - Time: $\Theta(n^2)$
2. Better solution: Use sort (see hw)
 - Time: $\Theta(???)$
3. Best: 1) Insert elements of A into hash table H ; 2) For each x in A , search for $t - x$.
 - Time: $\Theta(n)$

Further Immediate Application

- Spellchecker
- Database

Hash Table

Summary

- Choice of the hash function.
- Collision resolution scheme.
- Hash table size and rehashing.
- Time complexity of **hash table** versus **sorted array**
 - insert(): $O(1)$ versus $O(n)$
 - find(): $O(1)$ versus $O(\log n)$
- When **NOT** to use hash?
 - **Rank search**: return the k-th largest item.
 - **Sort**: return the values in order.